

ISA details

No special modifications have been made to the proposed ISA. If anything, the only point of note is that the write port for the register file has indeed been kept as the first operand of all the R-type instructions (as opposed to having ports #1 and #2 be the more conventional R_s or R_t , with the third port taking the place of R_d) and the corresponding M-type instructions in `ADDI` and `LD{B,W}`. Notice that, because of this terminological simplification, `R1` and `R3` must be muxed for instructions `ST{B,W}` and `BEQ` (notice the “s.t. *r1* is moved to *r3*” comment, see mux `R1R3SRC` within the diagram figure of the processor). The `SLT` instruction is a byproduct of the ALU block’s implementation and is rarely used, if at all, in the tested code and the performance tests. The `BEQ` instruction shares the same encoding as M-type instructions, while the `JMP` instruction does its own thing.

Instructions

| | | | | |
|------|---|------|-------------|---|
| 0x0 | R | ADD | r1, r2, r3 | // r1 <- r2 + r3 |
| 0x1 | R | SUB | r1, r2, r3 | // r1 <- r2 - r3 |
| 0x2 | R | AND | r1, r2, r3 | // r1 <- r2 AND r3 |
| 0x3 | R | OR | r1, r2, r3 | // r1 <- r2 OR r3 |
| 0x4 | R | SLT | r1, r2, r3 | // r1 <- "1" if r2 < r3, otherwise r1 <- "0" |
| 0x5 | R | MUL | r1, r2, r3 | // r1 <- r2 * r3 |
| 0x7 | M | ADDI | r1, r2, imm | // r1 <- r2 + imm |
| 0x10 | M | LDB | r1, r2, imm | // r1 <- mem[r2 + imm] |
| 0x11 | M | LDW | r1, r2, imm | // r1 <- mem[r2 + imm] |
| 0x12 | M | STB | r1, r2, imm | // mem[r2 + imm] <- r1, s.t. r1 is moved to r3 |
| 0x13 | M | STW | r1, r2, imm | // mem[r2 + imm] <- r1, s.t. r1 is moved to r3 |
| 0x30 | B | BEQ | r1, r2, imm | |
| | | | | // if r1 == r2, PC = PC_plus4 + imm, s.t. r1 is moved to r3 |
| 0x31 | B | JMP | imm | |
| | | | | // PC = concat(PC(31 downto 27, imm, "00"), imm is instr(24 downto 0) |

Encoding details

[R-type]

| | | | | |
|--------|-------|-------|-------|--------|
| 31-25 | 24-20 | 19-15 | 14-10 | 9-0 |
| opcode | dst | src1 | src2 | zeroes |

[M-type]

| | | | |
|--------|-------|-------|--------|
| 31-25 | 24-20 | 19-15 | 14-0 |
| opcode | dst | src1 | offset |

[BEQ-type] opcode_31-25 dst_24-20 src1_19-15 offset_14-0

[JMP-type] opcode_31-25 offset_24-0

Sanity-check program and n-th Fibonacci number calculation

The following programs have been used to verify the correct functionality of the processor in its two stages (single cycle and pipelined). Both of them have been pulled from the [CMOS VLSI Design: A Circuits and Systems Perspective - 4th Edition - Weste, Harris] book, where an introductory section of the book (more specifically, chapter 1.7) uses a simple multicycle MIPS to explore the concept “*hierarchical design*”, as well as making a demonstration of what its actual synthesised, physical implementation looks like (refer to chapters 1.9 and 1.10, “*Circuit design*” and “*Physical design*” for further information). Along with a high-level block diagram explaining how the processor is interconnected come the two programs in question: a simple instruction checker and a n-th Fibonacci number calculator. Figures from both the book and the used assembly code are included. The Fibonacci program is also used as a performance test comparing the single cycle and the pipelined versions of the processor. The actual assembly code that is loaded into the memory .dat files resides within the `asmtoisa.py` script.

Book figure 1

```
# mipstest.asm
# 9/16/03 David Harris David_Harris@hmc.edu
#
# Test MIPS instructions. Assumes little-endian memory was
# initialized as:
# word 16: 3
# word 17: 5
# word 18: 12

main:      #Assembly Code      effect      Machine Code
          lb $2, 68($0)        # initialize $2 = 5      80020044
          lb $7, 64($0)        # initialize $7 = 3      80070040
          lb $3, 69($7)        # initialize $3 = 12     80e30045
          or $4, $7, $2        # $4 <= 3 or 5 = 7      00e22025
          and $5, $3, $4       # $5 <= 12 and 7 = 4     00642824
          add $5, $5, $4       # $5 <= 4 + 7 = 11      00a42820
          beq $5, $7, end      # shouldn't be taken     10a70008
          slt $6, $3, $4       # $6 <= 12 < 7 = 0      0064302a
          beq $6, $0, around   # should be taken        10c00001
          lb $5, 0($0)         # shouldn't happen       80050000
around:    slt $6, $7, $2       # $6 <= 3 < 5 = 1       00e2302a
          add $7, $6, $5       # $7 <= 1 + 11 = 12     00c53820
          sub $7, $7, $2       # $7 <= 12 - 5 = 7      00e23822
          j end                # should be taken        0800000f
          lb $7, 0($0)         # shouldn't happen       80070000
end:       sb $7, 71($2)       # write adr 76 <= 7     a0470047
          .dw 3                00000003
          .dw 5                00000005
          .dw 12               0000000c
```

Extracted from section A.12.1 of the book, final output of the program should place a 7 in address 76

Book figure 2

```
int fib(void)
{
    int n = 8;           /* compute nth Fibonacci number */
    int f1 = 1, f2 = -1; /* last two Fibonacci numbers */

    while (n != 0) {      /* count down to n = 0 */
        f1 = f1 + f2;
        f2 = f1 - f2;
        n = n - 1;
    }
    return f1;
}
```

FIGURE 1.50 C Code for Fibonacci program

Book figure 3

```
# fib.asm
# Register usage: $3: n $4: f1 $5: f2
# return value written to address 255
fib:  addi $3, $0, 8      # initialize n=8
      addi $4, $0, 1      # initialize f1 = 1
      addi $5, $0, -1     # initialize f2 = -1
loop: beq $3, $0, end     # Done with loop if n = 0
      add $4, $4, $5      # f1 = f1 + f2
      sub $5, $4, $5      # f2 = f1 - f2
      addi $3, $3, -1     # n = n - 1
      j loop              # repeat until done
end:  sb $4, 255($0)      # store result in address 255
```

FIGURE 1.51 Assembly language code for Fibonacci program

Code figure 1

| main: | #Assembly Code | effect | Machine Code | Data memory | Bytenum |
|---------|--------------------|---------------------------|--------------|-------------|---------|
| | LDB 2 0 68 | # initialize \$2 = 5 | 20200044 | 80020044 | 0 |
| | LDB 7 0 64 | # initialize \$7 = 3 | 20700040 | 80070040 | 4 |
| | LDB 3 7 69 | # initialize \$3 = 12 | 20338045 | 80e30045 | 8 |
| | OR 4 7 2 | # \$4 <= 3 or 5 = 7 | 06438800 | 00e22025 | 12 |
| | AND 5 3 4 | # \$5 <= 12 and 7 = 4 | 04519000 | 00642824 | 16 |
| | ADD 5 5 4 | # \$5 <= 4 + 7 = 11 | 00529000 | 00a42820 | 20 |
| | BEQ 5 7 8 (end) | # shouldn't be taken | 60538008 | 10a70008 | 24 |
| | SLT 6 3 4 | # \$6 <= 12 < 7 = 0 | 08619000 | 0064302a | 28 |
| | BEQ 6 0 1 (around) | # should be taken | 60600001 | 10c00001 | 32 |
| | LDB 5 0 0 | # shouldn't happen | 20500000 | 80050000 | 36 |
| around: | SLT 6 7 2 | # \$6 <= 3 < 5 = 1 | 08638800 | 00e3202a | 40 |
| | ADD 7 6 5 | # \$7 <= 1 + 11 = 12 | 00731400 | 00c53820 | 44 |
| | SUB 7 7 2 | # \$7 <= 12 - 5 = 7 | 02738800 | 00e23822 | 48 |
| | JMP 15 | # should be taken | 6200000f | 0800000f | 52 |
| | LDB 7 0 0 | # shouldn't happen | 20700000 | 80070000 | 56 |
| end: | STB 7 2 71 | # write adr 76 <= \$7 | 24710047 | a0470047 | 60 |
| | STW 7 0 0 | # mem(0) <= \$7 word-wise | 26700000 | 00000003 | 64 |
| | STB 7 0 1 | # mem(1) <= \$7 byte-wise | 24700001 | 00000005 | 68 |
| | LDB 1 0 1 | # \$1 <= mem(1) byte-wise | 20100001 | 0000000c | 72 |
| | LDW 1 0 0 | # \$1 <= mem(0) word-wise | 22100000 | | 76 |
| | ADDI 3 3 -1 | # \$3 <= \$3 + -1 | 0e31ffff | | 80 |
| | MUL 3 3 3 | # \$3 <= \$3 * \$3 | 0a318c00 | | 84 |

Code figure 2

```

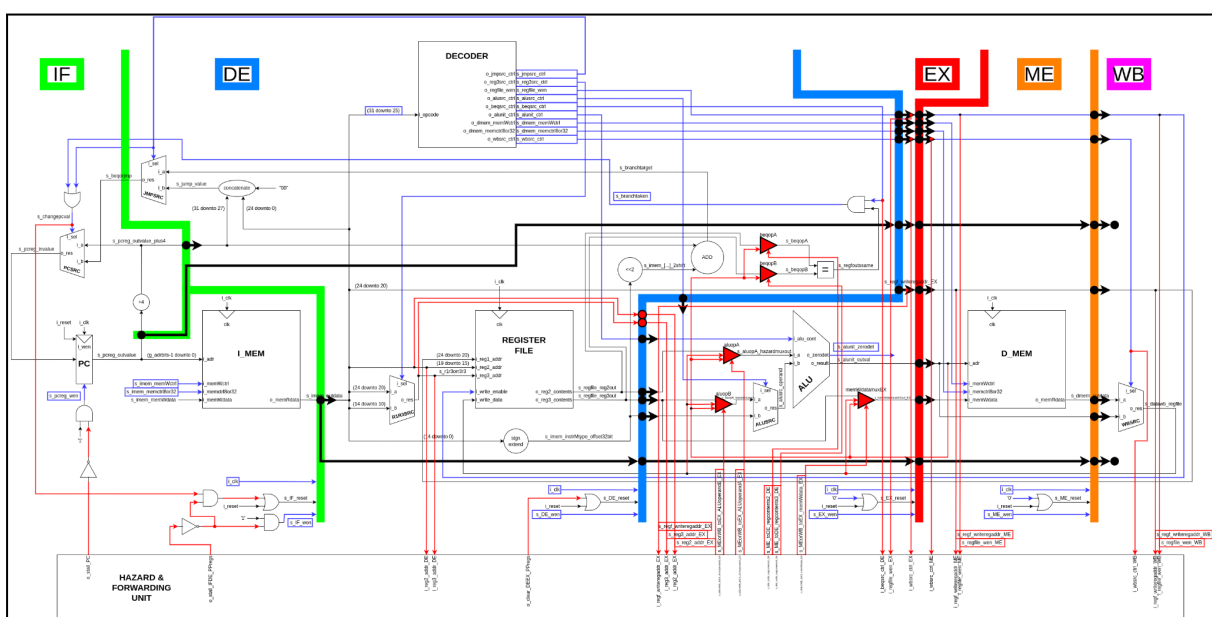
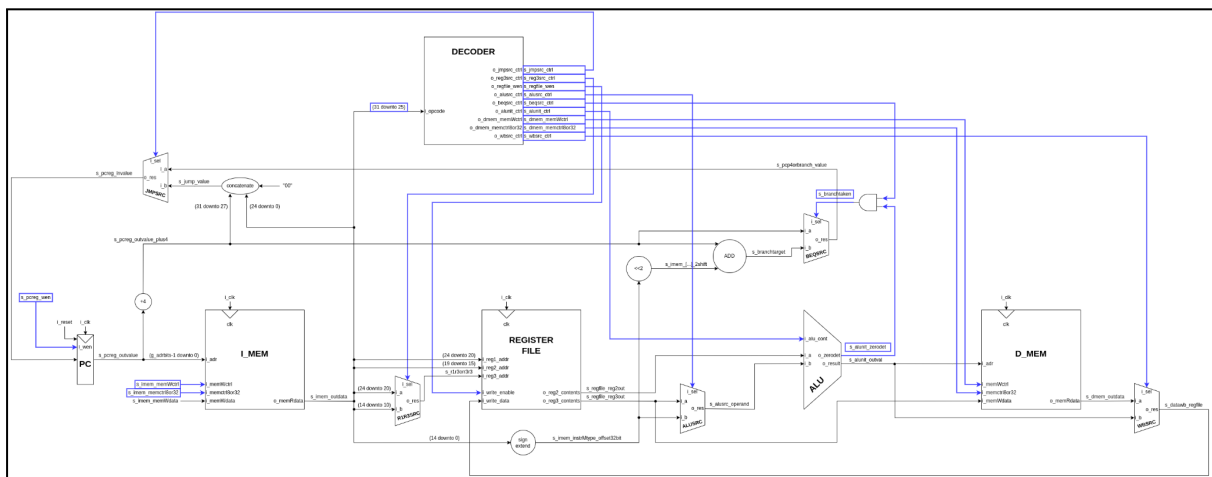
# fib.asm
# Register usage: $3: n $4: f1 $5: f2
# return value written to address 255
fib:  addi $3, $0, 8      # initialize n=8
      addi $4, $0, 1      # initialize f1 = 1
      addi $5, $0, -1     # initialize f2 = -1
loop: beq $3, $0, end     # Done with loop if n = 0
      add $4, $4, $5      # f1 = f1 + f2
      sub $5, $4, $5      # f2 = f1 - f2
      addi $3, $3, -1     # n = n - 1
      j loop              # repeat until done
end:  sb $4, 255($0)      # store result in address 255

#####
# translation to "bespoke" format
#####
fib:  ADDI 3 0 8          # initialize n=8
      ADDI 4 0 1          # initialize f1 = 1
      ADDI 5 0 -1         # initialize f2 = -1
loop: BEQ 3 0 4           # Done with loop if n = 0
      ADD 4 4 5           # f1 = f1 + f2
      SUB 5 4 5           # f2 = f1 - f2
      ADDI 3 3 -1         # n = n - 1
      JMP 3               # repeat until done
end:  STB 4 0 255         # store result in address 255

```

Processor diagrams

All the signals names correspond to the actual implementation in code. See the .drawio file for further reference. The single cycle implementation has been used as a baseline to confirm the results of the pipelined version and a comparison of both regarding the performance tests is provided. The relatively low level of idiosyncrasies and the (much to my chagrin) simplicity of the pipeline design lends itself to being more easily explained via visual queues, or at least that would be my take on the matter. As such, forgoing a per-instruction-type text introduction to what the processor does for each instruction and how they relate to each other in the pipeline, the .drawio file provided with the project files lays out a step-by-step extension-process of how to go from the single-cycle version of the processor to the fully functional pipelined version. Each tab of the .drawio file is aptly named to refer to the motivation and purpose behind the extension (note: Ctrl+Shift+H to centre the design in the window and thus see each transition from a sort of “framed” fixed position). The single-cycle version corresponds to the datapath.vhd file (simulated with the datapath_tb.vhd file) and the pipelined version resides in the datapath_PP{, _tb}.vhd files.



Performance test results

Four performance tests are considered in this section: the n-th Fibonacci calculator, the proposed buffer-sum process, the proposed mem_copy process and the proposed matrix multiplication process. Below are the figures for the original code and their implementation in assembly.

Fibonacci C code

```
int fib(void)
{
    int n = 8;           /* compute nth Fibonacci number */
    int f1 = 1, f2 = -1; /* last two Fibonacci numbers */

    while (n != 0) {     /* count down to n = 0 */
        f1 = f1 + f2;
        f2 = f1 - f2;
        n = n - 1;
    }
    return f1;
}
```

FIGURE 1.50 C Code for Fibonacci program

Fibonacci code in its original (from the book) assembly format and the processor-adapted assembly code for our machine (below)

```
# fib.asm
# Register usage: $3: n $4: f1 $5: f2
# return value written to address 255
fib:  addi $3, $0, 8      # initialize n=8
      addi $4, $0, 1      # initialize f1 = 1
      addi $5, $0, -1     # initialize f2 = -1
loop: beq $3, $0, end     # Done with loop if n = 0
      add $4, $4, $5      # f1 = f1 + f2
      sub $5, $4, $5      # f2 = f1 - f2
      addi $3, $3, -1     # n = n - 1
      j loop              # repeat until done
end:  sb $4, 255($0)      # store result in address 255

#####
# translation to "bespoke" format
#####
fib:  ADDI 3 0 8          # initialize n=8
      ADDI 4 0 1          # initialize f1 = 1
      ADDI 5 0 -1         # initialize f2 = -1
loop: BEQ 3 0 4           # Done with loop if n = 0
      ADD 4 4 5           # f1 = f1 + f2
      SUB 5 4 5           # f2 = f1 - f2
      ADDI 3 3 -1         # n = n - 1
      JMP 3               # repeat until done
end:  STB 4 0 255         # store result in address 255
```

Actual assembly code in the asmt0isa.py script

```
· · #####  
· · # NeilWeste nth Fibonacci number  
· · #####  
· · "ADDI 3 0 9",  
· · "ADDI 4 0 1",  
· · "ADDI 5 0 -1",  
· · "BEQ 3 0 4",  
· · "ADD 4 4 5",  
· · "SUB 5 4 5",  
· · "ADDI 3 3 -1",  
· · "JMP 3",  
· · "STW 4 0 255",
```

Buffer-sum original code

```
• Buffer_sum  
  int a[128], sum = 0;  
  for (i=0; i<128; i++) { sum += a[i]; }
```

Buffer-sum assembly implementation

```
· · # perftest 1  
· · "ADDI 1 0 128", # limit of 128  
· · "ADD 2 0 0", · · · # sum  
· · "ADD 3 0 0", · · · # i  
· · "BEQ 1 3 4", · · · # start loop  
· · "LDB 4 3 0",  
· · "ADD 2 2 4", · · · # sum = a[i]  
· · "ADDI 3 3 1", · · · # i++  
· · "JMP 3", · · · · · # try loop condition again  
· · "ADDI 5 0 1", · · · # finish flag
```

Mem-copy original code

```
• Mem_copy
  int a[128], b[128];
  for (i=0; i<128; i++) { a[i] = 5; }
  for (i=0; i<128; i++) { b[i] = a[i]; }
```

Mem-copy assembly implementation

```
· # perftest 2
· "ADDI 1 0 128", # limit of 128
· "ADDI 2 0 5", # 5 value
· "ADD 3 0 0", # i
· "ADDI 4 0 0", # base pointer for a[128]
· "ADDI 5 0 128", # base pointer for b[128]
· "BEQ 1 3 4", # start loop #1
· "STB 2 4 0", # a[i] = 5
· "ADDI 4 4 1", # update pointer of a[128]
· "ADDI 3 3 1", # i++
· "JMP 5", # try loop condition again
· "SUB 4 4 3", # reset pointer of a[128]
· "ADDI 3 0 0", # reset i
· "BEQ 1 3 6", # start loop #2
· "LDB 6 4 0", # intermediate register to hold a[i] value
· "STB 6 5 0", # b[i] = a[i]
· "ADDI 4 4 1", # update pointer of a[128]
· "ADDI 5 5 1", # update pointer of b[128]
· "ADDI 3 3 1", # i++
· "JMP 12", # try loop condition again
· "ADDI 31 0 1", # finish flag
```


Matrix multiply original code

- Matrix multiply

```
int a[128][128], b[128][128], c[128][128];
for (i=0; i<128; i++) {
    for(j =0;j<128;j++) {
        c[i][j] = 0;
        for(k = 0; k <128; k++ ) {
            c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
    }
}
```

Matrix multiplication assembly implementation

```
# perftest 3
"ADDI 1 0 10", # limit of 128
"ADD 2 0 0", # i
"ADD 3 0 0", # j
"ADD 4 0 0", # k
"ADDI 5 0 0", # base pointer for a[][]
"ADDI 6 0 0", # base pointer for b[][]
"ADDI 7 0 0", # base pointer for c[][]
"BEQ 1 2 22", # start loop #1
"BEQ 1 3 18", # start loop #2
"MUL 10 2 1", # i*128
"ADD 11 10 3", # (i*128)+j
"ADD 12 11 7", # (i*128)+j+base_c_pointer
"STB 0 12 0", # c[i][j] = 0
"BEQ 1 4 10", # start loop #3
# a[i][k] section
"MUL 13 2 1", # i*128
"ADD 14 13 4", # (i*128)+k
"ADD 15 14 5", # (i*128)+k+base_a_pointer
# b[k][j] section
"MUL 16 2 1", # k*128
"ADD 17 16 3", # (k*128)+j
"ADD 18 17 6", # (k*128)+j+base_b_pointer
# a[i][k]*b[k][j] section
"MUL 19 15 18",
# c[i][j]=c[i][j]+a[i][k]*b[k][j] section
"ADD 12 12 19",
"ADDI 4 4 1", # k++
"JMP 13", # try loop #3 condition again
"ADDI 4 0 0", # reset k
"ADDI 3 3 1", # j++
"JMP 8", # try loop #2 condition again
"ADDI 3 0 0", # reset j
"ADDI 2 2 1", # i++
"JMP 7", # try loop #1 condition again
"ADDI 2 0 0", # reset i
"ADDI 8 0 1", # finish flag
```

Results:

| | expected value (single cycle output) | single-cycle #cycles | pipelined #cycles | single-to-pp ratio | pp ipc |
|-----------------|--|-------------------------|----------------------|-----------------------|--------------|
| fib 10 | 34 | 551 | 691 | 1.254083485 | 0.7973950796 |
| fib 50 | 3483774753 | 2551 | 3091 | 1.211681693 | 0.8252992559 |
| fib 100 | 3405478146 | 5051 | 6091 | 1.205899822 | 0.8292562798 |
| fib 200 | 305885837 | 10051 | 12091 | 1.202964879 | 0.8312794641 |
| fib 1000 | 4057268386 | 50051 | 60091 | 1.200595393 | 0.8329200712 |

| | single-cycle #cycles | pipelined #cycles | single-to-pp ratio | pp ipc |
|---------------------------|-------------------------|-------------------|--------------------|--------------|
| buffer-sum | 6451 | 9071 | 1.406138583 | 0.7111674567 |
| mem-copy | 15461 | 19371 | 1.252894379 | 0.7981518765 |
| mx-multiply (i=10) | 119601 | 143051 | 1.196068595 | 0.8360724497 |